

# Building Docker Images

## What is a Dockerfile? How do you build a Docker image? #

**Dockerfile:** A file containing a series of instructions on how to build a Docker image

### Example:

- Dockerfile

```
# Use an official Python runtime as the base image
FROM python:3.8-slim

# Sets the working directory inside the container
# Conceptually similar to `cd` command
WORKDIR /app

# Copies files/directories from the host to the container
COPY . /app

# Install any needed packages specified in requirements.txt
RUN pip install --no-cache-dir -r requirements.txt

# Exposes a port for external access
# Make port 80 available to the world outside this container
EXPOSE 80

# Define the command to run the application
CMD ["python", "app.py"]
```



- Build the image using:

```
docker build -t my-python-app .
```



- Create container from the image:

## On this page

What is a Dockerfile? How do you build a Docker image?

What is a Base Image?

Why is the Base Image a Key Choice?

What is the difference between ENTRYPPOINT and CMD in a Dockerfile? explain with example

Difference between ADD and COPY in a Dockerfile

How do you tag a Container image? Why is tagging important?

Explain the layered approach to building a Container image?

How can you design Dockerfiles to maximize layer reuse?

What is a multi-stage build, and how can it be used to reduce the size of a container image?

Is it possible to make changes to an existing Container image?

What is the purpose of the .dockerignore file?

What Are Alternatives to Writing Dockerfile?

What is OCI and Why Is It Important?

```
docker run -d -p 80:80 my-python-app
```

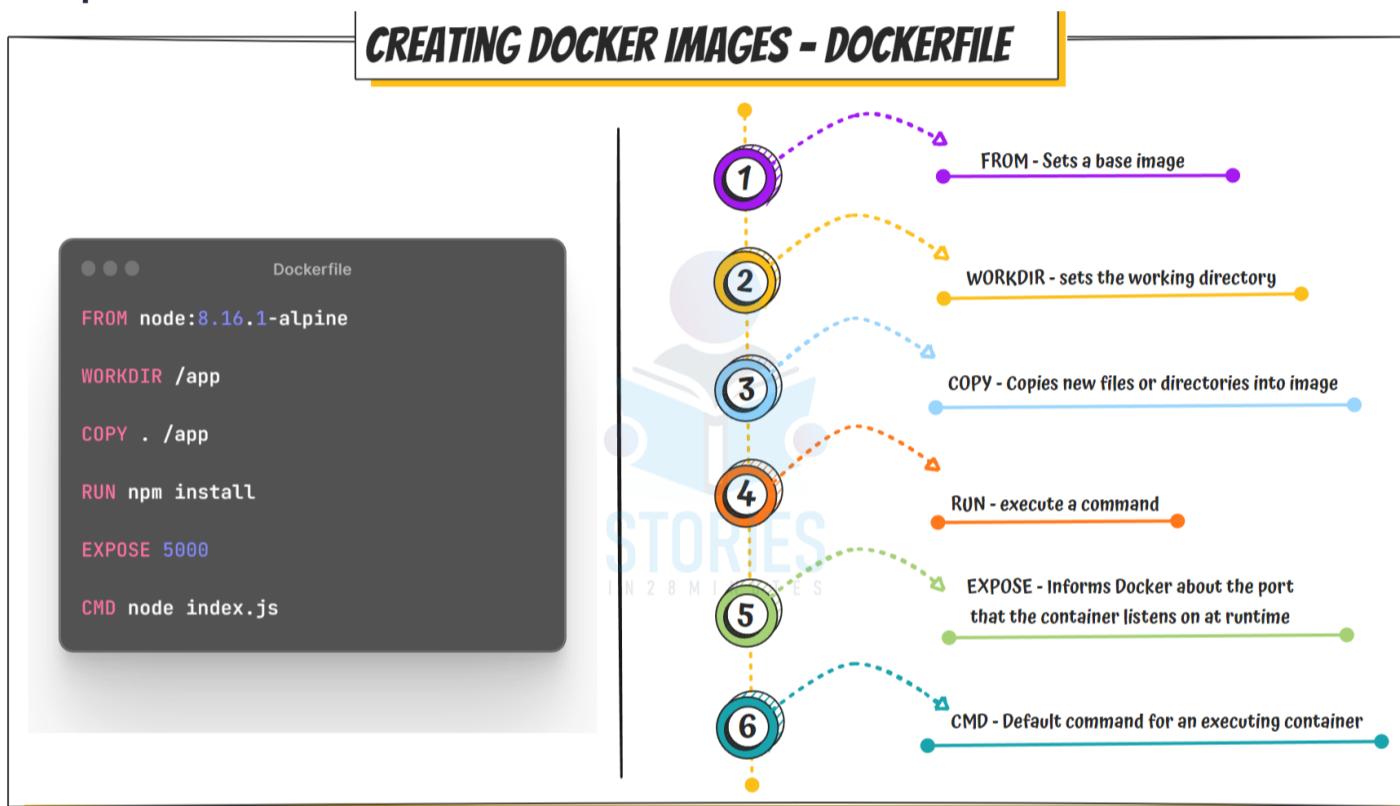
```
# Test  
curl http://localhost
```



**Explore** the following projects with detailed guides:

- [hello-world-java](#)
- [hello-world-node](#)
- [hello-world-python](#)

**Example:**



## What is a Base Image? #

- **Starting Point for Builds:** Provides an initial starting point to build your container images
- **Add Custom Things:** Add your application code, dependencies, and configurations
- **Huge Variety:** You have a variety of Base Images
  - **OS Images:** Examples include ubuntu:22.04, debian:12, alpine:3.20;
  - **Language Runtimes:** Examples include openjdk:8-jdk-slim, node:30-alpine, python:4.12-slim; these come with pre-installed runtime environments
  - **Minimal & Secure Options:** Such as alpine:3.20 or gcr.io/distroless/java21; they are tiny and have a smaller attack surface
  - **Blank Slate:** The 'scratch' image gives a completely empty starting point for maximum control

## Why is the Base Image a Key Choice? #

- **Size:** Smaller images lead to faster downloads and less storage usage
- **Security:** Choosing a minimal and well-maintained image reduces vulnerability risks
- **Performance:** Influences startup time and runtime efficiency of your container
- **Community Support:** Popular images often have better documentation and troubleshooting resources

## What is the difference between ENTRYPOINT and CMD in a Dockerfile? explain with example #

- **ENTRYPOINT**: Defines the main executable
  - **Main executable:** Sets the container's primary command
  - **Not easily overridden:** Not overridden by arguments unless `--entrypoint` is used.
- **CMD**: Default arguments to the ENTRYPOINT
  - **Default execution:** Provides default arguments for container run
  - **Can be easily overridden:** Can be overridden by arguments passed to `docker run`.

### ENTRYPOINT Example:

```
## Use an official Python runtime as a parent image
FROM python:3.8-slim

# Set the working directory
WORKDIR /app

# Copy the current directory contents into the container
COPY . /app

# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt

# Set the entrypoint to run the Python interpreter with app.py
ENTRYPOINT ["python", "app.py"]
```



- **ENTRYPOINT** instruction sets the entry point for the container as `python app.py`
- When the container starts, it will run `python app.py` by default, and you cannot change this behavior by passing arguments to `docker run`

- To override the `ENTRYPOINT`, you need to use the `--entrypoint` flag with `docker run`:

```
docker run --entrypoint /bin/bash myimage
```



#### `ENTRYPOINT + CMD`:

```
# Use an official Python runtime as a parent image
FROM python:3.8-slim

# Set the working directory
WORKDIR /app

# Copy the current directory contents into the container
COPY . /app

# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt

# Set the entrypoint to run the Python interpreter
ENTRYPOINT ["python"]

# Define the default argument to run
CMD ["app.py"]
```



- This Dockerfile sets `ENTRYPOINT` to `python` and `CMD` to `app.py`
  - When the container starts, it will run `python app.py`
  - You can override the `CMD` argument like this:

```
docker run myimage another_script.py
```



#### `CMD Alone Example`:

```
# Use an official Python runtime as a parent image
FROM python:3.8-slim

# Set the working directory
WORKDIR /app

# Copy the current directory contents into the container
COPY . /app

# Install dependencies
```

```
RUN pip install --no-cache-dir -r requirements.txt

# Define the default command to run when the container starts
CMD ["python", "app.py"]
```



- The `CMD` instruction specifies that the container should run `python app.py` by default
- You can override this command by specifying a different command when you run the container

```
docker run myimage python another_script.py
```



## Difference between ADD and COPY in a Dockerfile #

- `COPY`: Simply copies files and directories from the host to the container
- `ADD`: More powerful. Can copy files and directories and also supports URL downloads and automatic extraction of compressed files

**Example:**

```
# Copy a local file into the /app directory in the image
COPY localfile.txt /app/

# Download a file from a URL
# and unzip it into the /app directory in the image
ADD http://example.com/file.tar.gz /app/
```



## How do you tag a Container image? Why is tagging important? #

- **Tagging Command:** Use `docker build -t imagename:tag .` to assign a tag
- **Importance:**
  - **Version Control:** Helps identify image versions
    - `microservice-a:v1, microservice-a:v2, microservice-a:v3`
    - `mysql:5.7, mysql:6.0`
    - `openjdk:8.0.0, openjdk:25.0.0`
  - **Rollbacks:** Easily revert to older versions of your image if newer ones fail

- **Flexibility:** Gives teams flexibility w.r.t. deploying different versions in different environments (dev, staging, prod).
- **Multiple Tags:** You can have multiple tags for an image
  - `latest` to indicate the latest version. Whoever want to use the latest version always can use the latest tag. `myapp:latest`

### Example:

```
# Build an myapp image with a tag 1.0.0
docker build -t myapp:1.0.0 .

# You can have multiple tags for an image
# Add additional tag latest to myapp image with tag 1.0.0
docker tag myapp:1.0.0 myapp:latest
```



## Explain the layered approach to building a Container image? #

- **Layered builds:** Container images are created in layers
- **Example:** Most instructions in a Dockerfile, such as `FROM`, `RUN`, `COPY`, create a new, read-only layer on top of the previous one.
- **Why is this important?**
  - **Layers are cached:** Layers are cached for faster and more efficient builds
  - **However Be Cautious:** Changing one layer triggers a rebuild of subsequent layers
  - **Example Below:** In Changed Dockerfile, Every layer from `Instruction 2 CHANGED` is rebuilt again

### Example Dockerfile

```
Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
```



### Changed Dockerfile

```
Instruction 1
Instruction 2 CHANGED
Instruction 3
```



## How can you design Dockerfiles to maximize layer reuse? #

- **Application code changes often:** In applications, code changes often, libraries/dependencies change infrequently
  - Application code evolves to add new features, fix bugs, and adapt to changing requirements
  - Libraries and dependencies, on the other hand, tend to have more stable APIs and are updated less frequently, often for security patches, performance improvements, or new features
- **Recommended:** Build Layers with Libraries and Dependencies First and Code Later

### Non-Optimized Dockerfile:

```
# Use an official Node.js runtime as a parent image
FROM node:24

# Set the working directory
WORKDIR /usr/src/app

# Copy application code + dependencies files
COPY . .

# Install dependencies - TAKES TIME!
RUN npm install

# Expose the application port
EXPOSE 3000

# Run the application
CMD ["node", "app.js"]
```



- In this Dockerfile, the `COPY . .` command copies the entire application code (`app.js`, ...) + dependencies files (`package.json` and `package-lock.json`) before running `npm install`
- This means any change in the application code (even a small one) will cause Docker to invalidate the cache for the `COPY . .` instruction and all subsequent instructions, leading to frequent rebuilds of the `npm install` step

## Optimized Dockerfile

```
# Use an official Node.js runtime as a parent image
FROM node:24

# Set the working directory
WORKDIR /usr/src/app

# Copy only package.json and package-lock.json first
COPY package*.json ./

# Install dependencies - TAKES TIME!
RUN npm install

# Copy the rest of the application code
COPY . .

# Expose the application port
EXPOSE 3000

CMD ["node", "app.js"]
```



- In this optimized Dockerfile, `package.json` and `package-lock.json` files are copied first
- The `RUN npm install` step runs only if dependencies change (`package.json` or `package-lock.json`)
  - `RUN npm install` will NOT if ONLY application code (`app.js`, ...) changes!
- The `COPY . .` command copies the rest of the application code after installing dependencies.
- This ensures changes to the application code do not affect the dependencies layer, allowing Docker to cache the `npm install` step unless `package.json` or `package-lock.json` change

## What is a multi-stage build, and how can it be used to reduce the size of a container image? #

- **Scenario:** Imagine needing a full compiler or SDK to build your app — but once it's built, you only need the tiny binary to run it. Do you really want to ship all that extra baggage in production? There's a better way.
- **Multi-Stage Build:** A Docker technique that lets you use multiple images in one `Dockerfile` — one for building your app and another for running it
- **Build Big, Run Small:** Use a larger image (like `node`, `maven`, or `golang`) to compile your app — then copy just the final output into a much smaller runtime image

- **Reduces Image Size Dramatically:** Final image contains only what's necessary to run the app — no compilers, dev tools, or source code
- **Improves Security by Design:** Smaller images have fewer vulnerabilities because they include fewer packages and libraries
- **Makes CI/CD More Efficient:** Smaller images mean faster downloads, deployments, and less storage usage across environments
- **How Does It Work?:**
  - **Multiple FROM statements:** Each FROM starts a new build stage
  - **Artifact copying:** Transfer only necessary artifacts between stages to keep the final image small

**Golang Example 1: Larger Image Size** Final image includes unnecessary build tools, making it larger than needed

```
FROM golang:1.16

WORKDIR /app

COPY main.go .

RUN go build main.go

CMD ["./main"]
```



**Golang Example 1: Multi Stage Build** Separate build stage, copy only compiled binary to minimal base

```
# Stage 1: Build the Go application
FROM golang:1.16 AS builder

WORKDIR /app

COPY main.go .

RUN go build main.go

# Stage 2: Create a minimal image with the built binary
FROM alpine:latest # MINIMAL BASE IMAGE

WORKDIR /root/

COPY --from=builder /app/main . # COPY WHAT WE WANT!

CMD ["./main"]
```



# Is it possible to make changes to an existing Container image? #

- **Read-only design:** Container images are immutable — you cannot directly modify them.
  - This ensures consistency, reproducibility, and follows the principle of immutable infrastructure.
- **Commit changes:** However, Modifications can be made to a running container and then committed to create a new image
- **New image creation:** Instead of updating the existing image, a new one is created, following immutable infrastructure principles

**Example:** Commit changes from container & Create new image

- **Step 1:** Run a container from the `nginx:latest` image

```
docker run --name test1 -d -p 80:80 nginx

# Test
curl http://localhost:80
```



- **Step 2:** Modify the running container

```
docker exec test1 sh -c \
'echo "<h1 style="font-size:400px">Testing</h1>" > /usr/share/nginx/html/:

# Test
curl http://localhost:80
```



- **Step 3:** Create new container image from running container

```
docker commit test1 newnginx:modified
```



- **Step 4:** Run a container from the `newnginx:modified` image (newly created)

```
docker run --name test2 -d -p 81:80 newnginx:modified

# Test
curl http://localhost:81
```



## What is the purpose of the `.dockerignore` file? #

- **Exclude files:** `.dockerignore` prevents specific files and directories from being copied into the Container image during the build and make image lightweight
- **Pattern matching:** Define file and directory patterns (similar to `.gitignore`) that Docker should skip during the build process

Example `.dockerignore` File:

```
# Ignore node_modules directory
node_modules

# Ignore log files
*.log

# Ignore development or editor-specific files
.DS_Store
.vscode/
.idea/
```



## What Are Alternatives to Writing Dockerfile? #

- **Manual Dockerfile Challenges:** Writing good Dockerfile is challenging
  - Mistakes can lead to security issues, large images and long build times (inefficient caching)
- **Jib:** Java-specific tool that eliminates the need for Dockerfiles
  - **Layer Caching:** Built-in optimization to speed up image building process
  - **Security Features:** Reduce risks by automating best practices
- **Buildpacks:** Multi-language support (Java, Node.js, Python, etc.), no Dockerfile required
  - **Layer Caching:** Built-in optimization to speed up image building process
  - **Security Features:** Reduce risks by automating best practices
  - **Language Compatibility:** Supports a variety of programming languages

```
# Build a Java image with JIB
mvn compile com.google.cloud.tools:jib-maven-plugin:build \
  -Dimage=gcr.io/my-project/my-app

# Build a Java image with packer build pack
pack build my-spring-app \
  --builder paketobuildpacks/builder:base \
  --env BP_JVM_VERSION=24 \
```

```
--env BP_LOG_LEVEL=debug
```

```
# Use Spring Boot Maven Plugin (Uses packeto build pack)
mvn spring-boot:build-image
```



## What is OCI and Why Is It Important? #

- **Problem Before OCI:** Proprietary formats led to vendor lock-in and limited interoperability
- **Open Container Initiative(OCI):** Defines open, vendor-neutral specifications for containers
  - **Image Spec:** Standard format for container images
  - **Runtime Spec:** Standardization for running containers across different platforms
  - **Distribution Spec:** Standardized methods for pushing and pulling images through container registries (Supported by Docker Hub, Google Artifact Registry, Amazon Elastic Container Registry, ...)
- **Open Governance:** Managed by the Linux Foundation, ensuring open collaboration and transparency
- **OCI is Supported By Almost All Container Based Tools and Cloud Platforms:**  
Docker, Podman, Buildpacks, Kubernetes, containerd, AWS, Azure, Google Cloud

← Previous  
[Docker Fundamentals](#)

Next →  
[Managing Docker Containers](#)

### Keep Learning

[Home ↗](#)

[Springboot ↗](#)

[Cloud ↗](#)

### Our Products

[Roadmaps ↗](#)

[Flashcards ↗](#)

[Bookshelf ↗](#)